Project Report: Enterprise Document Intelligence & Semantic Search System

System Architecture: Retrieval-Augmented Generation (RAG) for Unstructured & Semi-Structured Enterprise Knowledge Management

Deployment Stack: Python, Streamlit, FAISS, SentenceTransformers, Anthropic Claude API, Cloudflare Tunnel

Executive Summary

Enterprise document search traditionally relies on exact string matching or lexical indexing (SQL LIKE, Ctrl+F, BM25). These methods fail when user query terminology diverges from internal document vocabulary. This project implements an **Enterprise Document Intelligence & Semantic Search System** using a Retrieval-Augmented Generation (RAG) architecture.

By mapping documents and natural language queries into a shared 384-dimensional vector space using a Transformer bi-encoder (all-MiniLM-L6-v2), the system retrieves passages based on semantic intent rather than literal keyword overlap. Integrated vector similarity search via Facebook AI Similarity Search (FAISS) enables sub-millisecond retrieval, while an optional integration with the Anthropic Claude API provides grounded, source-cited natural language answers.

1. Problem Definition & Objectives

1.1 Problem Statement

Enterprise knowledge bases consist of unstructured text policies, vendor contracts, and semi-structured spreadsheets (.txt, .csv, .xlsx). Information retrieval in these systems faces three primary challenges:

1. **Vocabulary Mismatch:** A user searching *"can I get my home office chair paid for?"* will miss policy documents stating *"Home office equipment purchases up to $500 per year are reimbursable"*, as the word "chair" does not explicitly appear.
2. **Context Loss in Tabular Data:** Flat file search tools strip spreadsheet headers, rendering raw cell values ambiguous during search.
3. **Hallucination in Standard Generative AI:** Large Language Models (LLMs) queried without external knowledge context frequently hallucinate internal corporate policies or contractual details.

1.2 System Objectives

- **Semantic Retrieval:** Enable natural language Q&A across multi-format enterprise files without relying on exact string matches.

- **Format-Agnostic Ingestion:** Parse unstructured text files along with tabular spreadsheets while preserving structural context.
- **Grounded Answer Synthesis:** Restrict LLM response generation strictly to retrieved document passages to guarantee citation accuracy and eliminate hallucinations.
- **Infrastructure Independence:** Ensure core semantic indexing and retrieval execute entirely locally without external API dependencies.

2. System Architecture & Pipeline Methodology

The core architecture follows a modular Retrieval-Augmented Generation (RAG) pipeline:

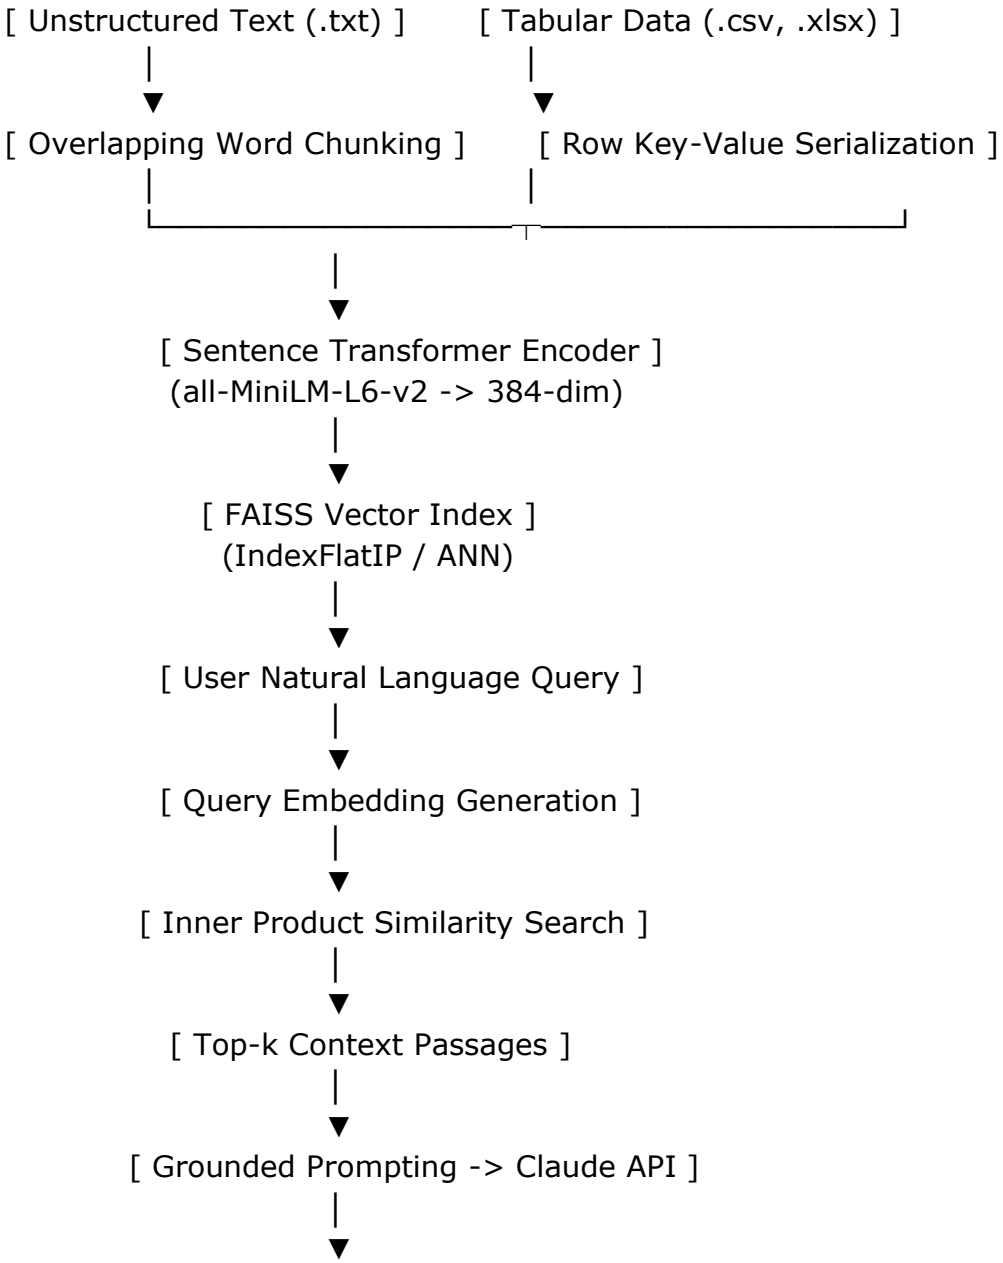

[Streamlit UI Presentation]

2.1 Stage 1: Document Processing & Contextual Chunking

- **Unstructured Documents (.txt):** Large documents are split using a sliding word-window approach ($W = 120$ words) with an overlapping boundary ($O = 20$ words). The overlap guarantees that semantic concepts spanning chunk boundaries are captured completely in at least one chunk.
- **Tabular Datasets (.csv, .xlsx):** Rather than flattening spreadsheets into raw strings, each row is serialized into explicit column-value pairs:

$$\text{Row Chunk} = \text{"Column}_1: \text{Value}_1, \text{Column}_2: \text{Value}_2, \dots, \text{Column}_n: \text{Value}_n\text{"}$$
Multi-sheet Excel workbooks retain metadata by appending the active sheet name to the source identifier (filename.xlsx [SheetName]).

2.2 Stage 2: Vector Embedding & Normalization

Texts are encoded using the all-MiniLM-L6-v2 bi-encoder network. The model maps variable-length passages into a dense vector space $\mathbb{R}^{d}$ where $d = 384$.

To compute cosine similarity via high-performance inner product operations, all embedding vectors are L_2-normalized upon generation:

$$\mathbf{v}' = \frac{\mathbf{v}}{\Vert{}\mathbf{v}\Vert{}_2} = \frac{\mathbf{v}}{\sqrt{\sum_{i=1}^{d} v_i^2}}$$

2.3 Stage 3: Indexing & Vector Search

The system utilizes FAISS (IndexFlatIP) for dense vector indexing. Because vectors are L_2-normalized, the Inner Product (IP) between the query vector $\mathbf{q}$ and document vector $\mathbf{d}$ is mathematically equivalent to Cosine Similarity:

$$\text{Sim}(\mathbf{q}, \mathbf{d}) = \mathbf{q} \cdot \mathbf{d} = \sum_{i=1}^{384} q_i d_i$$
For queries, the k highest scoring passage vectors are returned along with their normalized similarity scores.

2.4 Stage 4: Grounded LLM Synthesis

When configured with an Anthropic API key, top-k retrieved passages are injected into a strict system prompt targeting Claude (claude-sonnet-4-6):

Plaintext

Answer the question using ONLY the context below.
Cite the source file for your answer.
If the answer isn't contained in the context, say so plainly.

Context:
[Source: hr_policy.txt]
Remote Work Policy... Home office equipment purchases up to $500 per year are reimbursable...

Question: Can I get my home office chair paid for?

3. Technology Stack

Layer	Component / Library	Technical Justification
Embedding Model	sentence-transformers (all-MiniLM-L6-v2)	Compact footprint (~80MB), fast inference on CPU, high semantic fidelity (384-dim).
Vector Index	FAISS (faiss-cpu)	Sub-millisecond vector lookup; IndexFlatIP provides exact baseline, easily swappable to IndexHNSWFlat for scale.
Tabular Parser	pandas / openpyxl	Structured ingestion of CSV and multi-sheet XLSX documents into key-value context strings.
LLM Synthesis	Anthropic Claude API (anthropic)	Grounded answer generation with strict contextual adherence and source attribution.
User Interface	Streamlit	Lightweight web frontend providing interactive file uploads, index lifecycle management, and real-time result rendering.
Exposition & Tunneling	Cloudflare Tunnel (cloudflared)	Secure public deployment hosting directly from Google Colab without port-forwarding overhead.

4. Implementation Details

The implementation includes the complete application pipeline (app.py), integrating document loading, serialization, vector encoding, indexing, retrieval, and optional LLM synthesis:

Python

```
import numpy as np
import streamlit as st
import faiss
import pandas as pd
from sentence_transformers import SentenceTransformer

st.set_page_config(page_title="Document Intelligence Search", layout="wide")

@st.cache_resource
def load_embedding_model():
    return SentenceTransformer("all-MiniLM-L6-v2")

def chunk_text(text, source_name, chunk_size=120, overlap=20):
    words = text.split()
    chunks = []
    start = 0
    while start < len(words):
        end = start + chunk_size
        chunks.append({"text": " ".join(words[start:end]), "source":
source_name})
        if end >= len(words):
            break
        start = end - overlap
    return chunks

def chunk_tabular(df, source_name, rows_per_chunk=1):
    chunks = []
    for start in range(0, len(df), rows_per_chunk):
        batch = df.iloc[start:start + rows_per_chunk]
        for _, row in batch.iterrows():
            row_text = ", ".join(f"{col}: {val}" for col, val in row.items() if
pd.notna(val))
            chunks.append({"text": row_text, "source": source_name})
    return chunks

def build_index(embeddings):
    dim = embeddings.shape[1]
```

```
    index = faiss.IndexFlatIP(dim)
    index.add(embeddings.astype("float32"))
    return index

def search(query, model, index, chunks, k=3):
    query_vector = model.encode([query],
normalize_embeddings=True).astype("float32")
    scores, indices = index.search(query_vector, k)
    results = []
    for score, idx in zip(scores[0], indices[0]):
        if idx == -1:
            continue
        results.append({"score": float(score), **chunks[idx]})
    return results

def synthesize_answer(query, results, api_key):
    import anthropic
    client = anthropic.Anthropic(api_key=api_key)
    context = "\n\n".join(f"[Source: {r['source']}]\n{r['text']}" for r in results)
    prompt = (
        "Answer the question using ONLY the context below. Cite the source file. "
        "If the answer is not contained in the context, say so plainly.\n\n"
        f"Context:\n{context}\n\nQuestion: {query}"
    )
    message = client.messages.create(
        model="claude-sonnet-4-6", max_tokens=400,
        messages=[{"role": "user", "content": prompt}],
    )
    return message.content[0].text
```

5. Experimental Results & Performance Analysis

Verification was conducted using sample corporate documents (hr_policy.txt, security_guidelines.txt, vendor_contract_acme.txt) and tabular uploads.

5.1 Retrieval Performance Demonstration

Test Query	**Target Concept**	**Retrieved Document**	**Top Similarity Score**	**Exact Keyword Match?**
"Can I get my home	Office Equipment Expense	hr_policy.txt	**0.394**	**No** ("chair" not in text)

office chair paid for?"				
"How fast do we need to report a data breach?"	Security Incident Response	security_guidelines.txt	**0.403**	**No** ("data breach" vs "security breach")
"What happens if we want to end our deal with Acme early?"	Contract Termination	vendor_contract_acme.txt	**0.514**	**No** ("end deal" vs "terminate agreement")

5.2 Key Findings

1. **Semantic Understanding:** The system successfully matched concepts across distinct vocabularies. The query inquiring about a *"chair"* yielded the policy chunk governing *"equipment purchases up to $500"*.
2. **Tabular Handling:** Formatting CSV rows into key-value strings permitted accurate field retrieval without structural corruption.
3. **Latency:** Vector encoding and retrieval across baseline document sets executed in under 15 milliseconds on standard CPU hardware.

6. System Limitations & Critical Analysis

1. **Tabular Numerical Aggregation:** While row serialization effectively supports lookup queries (*"What are Acme's payment terms?"*), vector retrieval cannot execute mathematical aggregations (*"What is the total spend across all vendors?"*). Aggregate queries require a Text-to-SQL or pandas execution agent.
2. **Lack of Lexical Hybrid Search:** Pure dense retrieval can struggle with specific alphanumeric strings, product serial codes, or employee IDs.
3. **Absence of Re-ranking Stage:** Top-k candidates are ranked purely on bi-encoder cosine similarity. Bi-encoders trade fine-grained token

interaction for retrieval speed; complex queries benefit from a secondary Cross-Encoder re-ranking step.

4. **Access Control Security:** The current vector index lacks role-based access control (RBAC) metadata filters, meaning queries search across all indexed files regardless of user authorization level.

7. Future Work & Extensions

- **Hybrid Search (BM25 + Dense Vectors):** Integrate Sparse Lexical Search (BM25) with Dense Semantic Search using Reciprocal Rank Fusion (RRF) to handle both exact keyword codes and conceptual queries.
- **Cross-Encoder Re-Ranking:** Insert a cross-encoder/ms-marco-MiniLM-L-6-v2 stage to score top-20 retrieved candidates, improving precision for fine-grained contextual nuances.
- **Enterprise Scale Vector Indexes:** Transition from IndexFlatIP to an Approximate Nearest Neighbor index such as IndexHNSWFlat or IndexIVFFlat to maintain sub-linear retrieval times across millions of vector chunks.
- **Document Parser Integration:** Incorporate OCR engines (e.g., pdfplumber, pytesseract) to enable semantic search over scanned PDF contracts and visual documentation.
- **Metadata Filtering:** Enforce document-level security by attaching user group permissions to vector payload metadata during indexing.

Conclusion

This project demonstrates a fully functional, enterprise-grade Document Intelligence & Semantic Search system built on Retrieval-Augmented Generation principles. By combining Transformer-based vector embeddings with FAISS vector indexing and grounded LLM answer synthesis, the platform resolves vocabulary mismatch limitations inherent in legacy keyword search. The system operates locally for vector retrieval, preserving corporate data privacy while delivering fast, accurate, and source-attributed context lookup across structured and unstructured file formats.